


CSCD43: Database Systems Technology

Winter 2020
Prof. Nick Koudas



1



What Is a Database System?



- **Database:**
a very large, integrated collection of data.
- **Models a real-world enterprise**
 - Entities (e.g., teams, games)
 - Relationships
(e.g., The Raiders are playing in The Superbowl)
 - Also includes active components (e.g. “business logic”)
- A **Database Management System (DBMS)** is a software system designed to **store, manage, and facilitate access to databases**.

2



Is the WWW a DBMS?

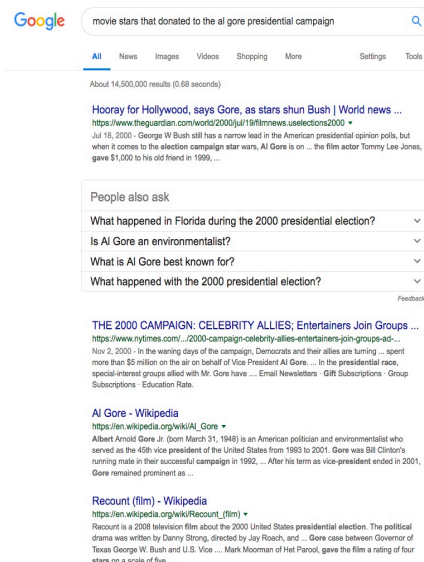
- Fairly sophisticated search available
 - crawler indexes pages on the web
 - Keyword-based search for pages
- **But, currently**
 - data is mostly unstructured and untyped
 - **search only:**
 - can't modify the data
 - can't get summaries, complex combinations of data
 - few guarantees provided for freshness of data, consistency across data items (data quality), fault tolerance, ...
 - Web sites (e.g. e-commerce) typically have a DBMS in the background to provide these functions.

3




“Search” vs. Query

- **What if you wanted to find out which movies stars donated to Al Gore's presidential campaign?**
- **Try “movie stars that donated to al gore presidential campaign” in your favorite search engine.**



4




“Search” vs. Query

- “Search” can return only what’s been “stored”
- E.g., best match Google, Bing top ten:



5



Is a File System a DBMS?

- **Thought Experiment 1:**
 - You and your project partner are editing the same file.
 - You both save it at the same time.
 - Whose changes survive?

A) Yours B) Partner’s C) Both D) Neither E) ???

- **Thought Experiment 2:**
 - You’re updating a file.
 - The power goes out.
 - Which of your changes survive?

Q: How do you write programs over a subsystem when it promises you only “???“ ?

A: Very, very carefully!!

A) All B) None C) All Since last save D) ???

6



Why Study Databases??



- The world is full of digital data, we better know
- how to manage them
- **Need for DBMS has exploded in the last years**
 - **Corporate**: retail swipe/clickstreams, “customer relationship mgmt”, “supply chain mgmt”, “data warehouses”, etc.
 - **Scientific**: digital libraries, Human Genome project, NASA Mission to Planet Earth, physical sensors, grid physics network
- **DBMS encompasses much of CS in a practical discipline**
 - OS, languages, theory, AI, multimedia, logic
 - Yet traditional focus on real-world apps

7



What's the intellectual content?



- **representing information**
 - data modeling
 - **languages and systems for querying data**
 - complex queries with real semantics*
 - over massive data sets
 - **Query processing**
 - Generating query answers from declarative statements
 - Query optimization
 - Query execution
 - **concurrency control for data manipulation**
 - controlling concurrent access
 - ensuring *transactional semantics*
 - **reliable data storage**
 - maintain data semantics even if you pull the plug
- * semantics: the meaning or relationship of meanings of a sign or set of signs

8



About the course: Workload

- **Projects with a “real world” focus:**
 - Modify the internals of a “real” open-source database system: PostgreSQL
 - Serious C system hacking
 - Measure the benefits of our changes
- **Exams – 1 Midterm & 1 Final**
- **Projects to be done in groups of up to 2**
 - Pick your partners ASAP

9



About the Course - Administrivia

- <http://www.cs.toronto.edu/~koudas>
 - Under teaching click on cscd43
- **Prof. Office Hours:**
 - Friday 11-12 in IC468 or by appointment
- **TAs: Yannis Xarchakos**
 - Office Hours: Tuesday after tutorial IC 400
 - Tutorials: Tuesday 9am-10am SW143
 - **Tutorials start Jan 21**

10



About the Course - Administrivia

- **Textbook**
 - Ramakrishnan et. al., Database Management Systems, 3rd edition
- **Cheating policy: zero tolerance**
- **Team Projects**
 - Three projects
 - Teams of up to 2
 - Submit; instructions on what exactly to submit will be provided with each project. Besides code, we will also be asking for performance evaluation for your algorithms compared with previous algorithms in the dbms.

11



- **A “free tasting” of things to come in this class:**
 - file systems & DBMSs
 - Query processing
 - Query Optimization
 - concurrent, fault-tolerant data management
 - DBMS architecture
- **Today’s lecture is from Chapter 1 in R&G**

12



OS Support for Data Management

- **Data can be stored in RAM**
 - this is what every programming language offers!
 - RAM is fast, and random access
 - Isn't this heaven?
- **Every OS includes a File System**
 - manages *files* on a magnetic disk or SSD
 - allows *open, read, seek, close* on a file
 - allows protections to be set on a file
 - drawbacks relative to RAM?

13



Database Management Systems

- **What more could we want than a file system?**
 - Simple, efficient *ad hoc*¹ queries
 - concurrency control
 - recovery
 - benefits of good data modeling
- **S.M.O.P.²? Not really...**
 - as we'll see this semester
 - in fact, the OS often gets in the way!

¹**ad hoc**: formed or used for specific or immediate problems or needs

²**SMOP**: Small Matter Of Programming

14



Describing Data: Data Models

- A ***data model*** is a collection of concepts for describing data.
- A ***schema*** is a description of a particular collection of data, using a given data model.
- The ***relational model of data*** is the most widely used model today.
 - Main concept: ***relation***, basically a table with rows and columns.
 - Every relation has a ***schema***, which describes the columns, or fields.

15



Example: University Database

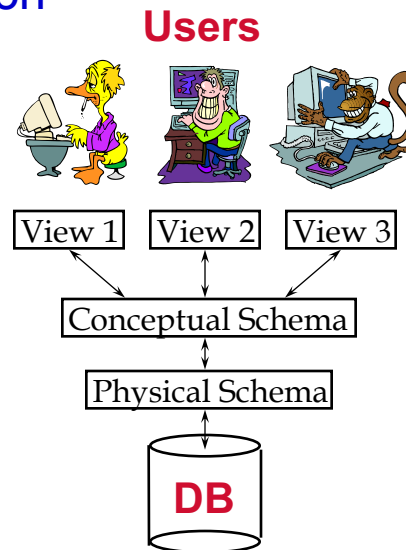
- **Conceptual schema:**
 - ***Students***(*sid: string, name: string, login: string, age: integer, gpa: real*)
 - ***Courses***(*cid: string, cname: string, credits: integer*)
 - ***Enrolled***(*sid: string, cid: string, grade: string*)
- **Physical schema:**
 - Relations stored as unordered files.
 - Index on first column of Students.
- **External Schema (View):**
 - ***Course_info***(*cid: string, enrollment: integer*)

16



Levels of Abstraction

- Views describe how users see the data.
- Conceptual schema defines logical structure
- Physical schema describes the files and indexes used.
- (sometimes called the ANSI/SPARC model)



17



Data Independence

- **Applications insulated from how data is structured and stored.**
- **Logical data independence:** Protection from changes in *logical* structure of data.
- **Physical data independence:** Protection from changes in *physical* structure of data.
- **Q: Why is this particularly important for DBMS?**

18



Concurrency Control

- **Concurrent execution of user programs: key to good DBMS performance.**
 - Disk accesses frequent, pretty slow
 - Keep the CPU working on several programs concurrently.
- **Interleaving actions of different programs: trouble!**
 - e.g., account-transfer & print statement at same time
- **DBMS ensures such problems don't arise.**
 - Users/programmers can pretend they are using a single-user system. (called “**Isolation**”)
 - Thank goodness! Don't have to program “very, very carefully”.

19



Transaction: An Execution of a DB Program

- **Key concept is a transaction: an **atomic sequence** of database actions (reads/writes).**
- **Each transaction, executed completely, must take the DB between **consistent** states.**
- **Users can specify simple integrity **constraints** on the data. The DBMS enforces these.**
 - Beyond this, the DBMS does not understand the semantics of the data.
 - Ensuring that a single transaction (run alone) preserves consistency is ultimately the user's responsibility!

20



Scheduling Concurrent Transactions

- **DBMS ensures that execution of $\{T_1, \dots, T_n\}$ is equivalent to some serial execution $T_1' \dots T_n'$.**
 - Before reading/writing an object, a transaction requests a lock on the object, and waits till the DBMS gives it the lock. All locks are held until the end of the transaction.
(Strict 2PL locking protocol.)
 - **Idea:** If an action of T_i (say, writing X) affects T_j (which perhaps reads X),
 - ... say T_i obtains the lock on X first
 - ... so T_j is forced to wait until T_i completes.This effectively orders the transactions.
 - What if
 - ... T_j already has a lock on Y
 - ... and T_i later requests a lock on Y ?(Deadlock!) T_i or T_j is aborted and restarted!


21



Ensuring Transaction Properties

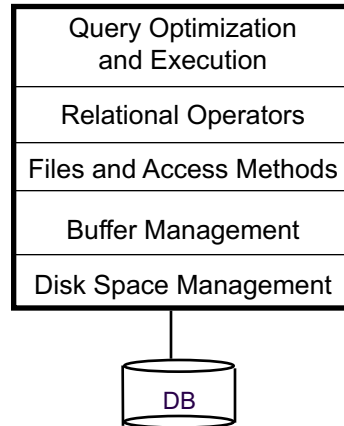
- **DBMS ensures atomicity (all-or-nothing property) even if system crashes in the middle of a Xact.**
- **DBMS ensures durability of committed Xacts even if system crashes.**
- **Idea: Keep a log (history) of all actions carried out by the DBMS while executing a set of Xacts:**
 - **Before** a change is made to the database, the corresponding log entry is forced to a safe location.
(WAL protocol; OS support for this is often inadequate.)
 - After a crash, the effects of partially executed transactions are undone using the log. Effects of committed transactions are redone using the log.
 - trickier than it sounds!

22



Structure of a DBMS

- A typical DBMS has a layered architecture.
- The figure does not show the concurrency control and recovery components.
- Each system has its own variations.
- You will see the “real deal” in PostgreSQL.
 - It’s a pretty full-featured example



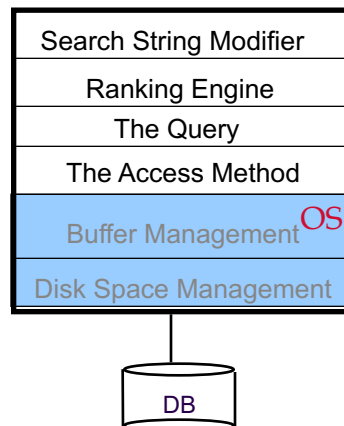
These layers must consider concurrency control and recovery

24



FYI: A text search engine

- **Less “system” than DBMS**
 - Uses OS files for storage
 - Just one access method
 - One hardwired query
 - regardless of search string
- **Typically no concurrency or recovery management**
 - Read-mostly
 - Batch-loaded, periodically
 - No updates to recover
 - OS a reasonable choice
- **Smarts: text tricks**
 - Search string modifier (e.g. “stemming” and synonyms)
 - Ranking Engine (sorting the output, e.g. by word or document popularity)
 - no semantics: WYGIWIGY



} Simple DBMS

→ There may be time to talk about some of these text tricks in this class, but it won’t be a focus.

25



Advantages of a DBMS

- **Data independence**
- **Efficient data access**
- **Data integrity & security**
- **Data administration**
- **Concurrent access, crash recovery**
- **Reduced application development time**
- **So why not use them always?**
 - Expensive/complicated to set up & maintain
 - This cost & complexity must be offset by need
 - General-purpose, not suited for special-purpose tasks (e.g. text search!)

26



Databases make these folks happy ...

- **DBMS vendors, programmers**
 - Oracle, IBM, MS, Spark, main memory DB vendors
- **End users in many fields**
 - Business, education, science, ...
- **DB application programmers**
 - Build enterprise applications on top of DBMSs
 - Build web services that run off DBMSs
- **Database administrators (DBAs)**
 - Design logical/physical schemas
 - Handle security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve



...must understand how a DBMS works

27



Summary (part 1)


- **DBMS used to maintain, query large datasets.**
 - can manipulate data and exploit *semantics*
- **Other benefits include:**
 - recovery from system crashes,
 - concurrent access,
 - quick application development,
 - data integrity and security.
- **Levels of abstraction provide data independence**
- **In this course we will explore:**
 - 1) How to be a sophisticated user of DBMS technology
 - 2) What goes on inside the DBMS

28



Summary, cont.

- DBAs, DB developers the bedrock of the information economy



- DBMS R&D represents a broad, fundamental branch of the science of computation

29




1. THE MAPREDUCE ABSTRACTION

30



The Map Reduce Abstraction for Distributed Algorithms



31



2. SPARK

32



Intro to Spark

- Spark is really a different implementation of the MapReduce programming model
- What makes Spark different is that it operates on Main Memory
- Spark: we write programs in terms of operations on resilient distributed datasets (RDDs).
- RDD (simple view): a collection of elements partitioned across the nodes of a cluster that can be operated on in parallel.
- RDD (complex view): RDD is an interface for data transformation, RDD refers to the data stored either in persisted store (HDFS) or in cache (memory, memory+disk, disk only) or in another RDD

33



RDDs in Spark


RDD: Resilient Distributed Datasets

- **Like a big list:**

- Collections of objects spread across a cluster, stored in RAM or on Disk

- **Built through parallel transformations**

- **Automatically rebuilt on failure**



Operations

- **Transformations** (e.g. map, filter, groupBy)
- **Make sure input/output match**