


Tree-Structured Indexes

Lecture 3
Chapter 10

"If I had eight hours to chop down a tree, I'd spend six sharpening my ax."

Abraham Lincoln



1



Introduction

- **Recall: 2 alternatives for data entries k^* :**
 - $\langle k, \text{rid of data record with search key value } k \rangle$
 - $\langle k, \text{list of rids of data records with search key } k \rangle$
- **Choice is orthogonal to the *indexing technique* used to locate data entries k^* .**
- **Tree-structured indexing techniques support both *range searches* and *equality searches*.**
- ***ISAM*: static structure; *B+ tree*: dynamic, adjusts gracefully under inserts and deletes.**
- **ISAM = ???**

Indexed Sequential Access Method

2



A Note of Caution

- **ISAM is an old-fashioned idea**
 - B+-trees are usually better, as we'll see
 - Though not *always*
- **But, it's a good place to start**
 - Simpler than B+-tree, but many of the same ideas
- **Upshot**
 - **Don't** brag about being an ISAM expert on your resume
 - **Do** understand how they work, and tradeoffs with B+-trees

3



Range Searches

- ``**Find all students with gpa > 3.0**'
 - If data is in sorted file, do binary search to find first such student, then scan to find others.
 - Cost of binary search can be quite high.
- **Simple idea: Create an 'index' file.**
 - Level of indirection again!




☛ *Can do binary search on (smaller) index file!*


4



ISAM



- Index file may still be quite large. But we can apply the idea repeatedly!




☛ Leaf pages contain *data entries*.

5



Example ISAM Tree

- Each node can hold 2 entries



6



Comments on ISAM

Data Pages

Index Pages

Overflow pages

- **File creation:** Leaf (data) pages allocated sequentially, sorted by search key. Then index pages allocated. Then space for overflow pages.
- **Index entries:** <search key value, page id>; they 'direct' search for *data entries*, which are in leaf pages.
- **Search:** Start at root; use key comparisons to go to leaf. $\text{Cost} \propto \log_F N$; $F = \# \text{ entries/index pg}$, $N = \# \text{ leaf pgs}$
- **Insert:** Find leaf where data entry belongs, put it there. (Could be on an overflow page).
- **Delete:** Find and remove from leaf; if empty overflow page, de-allocate.


➡ **Static tree structure:** *inserts/deletes affect only leaf pages.*

7



Example ISAM Tree


- Each node can hold 2 entries; no need for 'next-leaf-page' pointers. (Why?)



8




After Inserting 23*, 48*, 41*, 42* ...



9



... then Deleting 42*, 51*, 97*



➡ Note that 51 appears in index levels, but 51* not in leaf!

10



ISAM ---- Issues?


- **Pros**
 - ????
- **Cons**
 - ????

11



B+ Tree: The Most Widely Used Index

- Insert/delete at $\log_F N$ cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- Minimum 50% occupancy (except for root). Each node contains $d \leq \underline{m} \leq 2d$ entries. The parameter d is called the *order* of the tree.
- Supports equality and range-searches efficiently.




12



Example B+ Tree

- Search begins at root, and key comparisons direct it to a leaf (as in ISAM).
- Search for 5*, 15*, all data entries $\geq 24^*$...



➡ Based on the search for 15*, we know it is not in the tree!

13



B+ Trees in Practice

- **Typical order: 100. Typical fill-factor: 67%.**
 - average fanout = 133
- **Typical capacities:**
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- **Can often hold top levels in buffer pool:**
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes

14




Inserting a Data Entry into a B+ Tree

- **Find correct leaf L .**
- **Put data entry onto L .**
 - If L has enough space, *done!*
 - Else, must split L (into L and a new node $L2$)
 - Redistribute entries evenly, copy up middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- **This can happen recursively**
 - To split index node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- **Splits “grow” tree; root split increases height.**
 - Tree growth: gets wider or one level taller at top.

15




Example B+ Tree - Inserting 8*



16



Example B+ Tree - Inserting 8*





❖ Notice that root was split, leading to increase in height.

17



Inserting 8* into Example B+ Tree

- Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- Note difference between *copy-up* and *push-up*; be sure you understand the reasons for this.



18



Deleting a Data Entry from a B+ Tree


- **Start at root, find leaf L where entry belongs.**
- **Remove the entry.**
 - If L is at least half-full, *done!*
 - If L has only **$d-1$** entries,
 - Try to **re-distribute**, borrowing from *sibling* (adjacent node with same parent as L).
 - If re-distribution fails, **merge** L and sibling.
- **If merge occurred, must delete entry (pointing to L or sibling) from parent of L .**
- **Merge could propagate to root, decreasing height.**

19



Example Tree (including 8*)

Delete 19* and 20* ...




- **Deleting 19* is easy.**

20



Example Tree (including 8*)

Delete 19* and 20* ...




- Deleting 19* is easy.
- Deleting 20* is done with re-distribution.
Notice how middle key is *copied up*.

21



... And Then Deleting 24*

- Must merge.
- Observe *'toss'* of index entry (on right), and *'pull down'* of index entry (below).




22



Example of Non-leaf Re-distribution

- Tree is shown below *during deletion* of 24*. (What could be a possible initial tree?)
- In contrast to previous example, can re-distribute entry from left child of root to right child.




23



After Re-distribution

- Intuitively, entries are **re-distributed by 'pushing through'** the splitting entry in the parent node.
- It suffices to re-distribute index entry with key 20; we've re-distributed 17 as well for illustration.



24



Prefix Key Compression


- **Important to increase fan-out. (Why?)**
- **Key values in index entries only `direct traffic`; can often compress them.**
 - E.g., If we have adjacent index entries with search key values *Dannon Yogurt*, *David Smith* and *Devarakonda Murthy*, we can abbreviate *David Smith* to *Dav*. (The other keys can be compressed too ...)
 - Is this correct? Not quite! What if there is a data entry *Davey Jones*? (Can only compress *David Smith* to *Davi*)
 - In general, while compressing, must leave each index entry greater than every key value (in any subtree) to its left.
- **Insert/delete must be suitably modified.**

25



Bulk Loading of a B+ Tree

- **If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.**
 - Also leads to minimal leaf utilization --- why?
- **Bulk Loading can be done much more efficiently.**
- **Initialization: Sort all data entries, insert pointer to first (leaf) page in a new (root) page.**

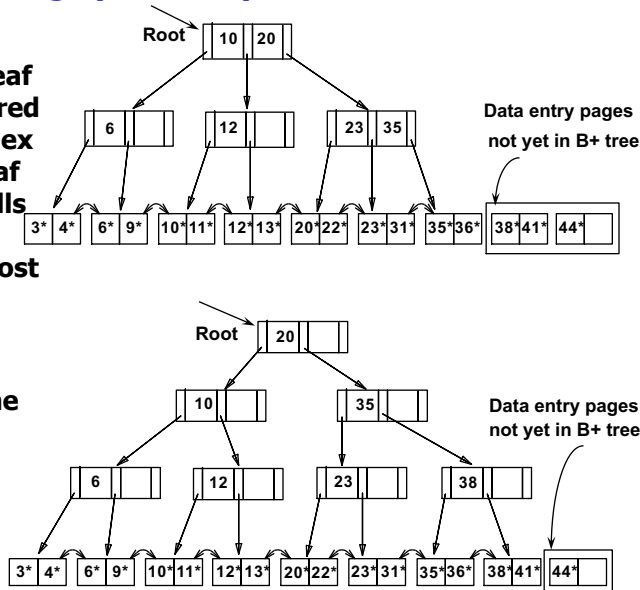


26



Bulk Loading (Contd.)

- Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits. (Split may go up right-most path to the root.)
- Much faster than repeated inserts, especially when one considers locking!



27



Summary of Bulk Loading

- **Option 1: multiple inserts.**
 - Slow.
 - Does not give sequential storage of leaves.
- **Option 2: Bulk Loading**
 - Has advantages for concurrency control.
 - Fewer I/Os during build.
 - Leaves will be stored sequentially (and linked, of course).
 - Can control “fill factor” on pages.

28



A Note on `Order`

- **Order (d) concept replaced by physical space criterion in practice (`at least half-full').**
 - Index pages can typically hold many more entries than leaf pages.
 - Variable sized records and search keys mean different nodes will contain different numbers of entries.
 - Even with fixed length fields, multiple records with the same search key value (*duplicates*) can lead to variable-sized data entries (if we use Alternative (2)).
- **Many real systems are even sloppier than this --- only reclaim space when a page is *completely* empty.**

29



Summary

- **Tree-structured indexes are ideal for range-searches, also good for equality searches.**
- **ISAM is a static structure.**
 - Only leaf pages modified; overflow pages needed.
 - Overflow chains can degrade performance unless size of data set and data distribution stay constant.
- **B+ tree is a dynamic structure.**
 - Inserts/deletes leave tree height-balanced; $\log_F N$ cost.
 - High fanout (**F**) means depth rarely more than 3 or 4.
 - Almost always better than maintaining a sorted file.

30



Summary (Contd.)

- Typically, 67% occupancy on average.
- Usually preferable to ISAM, adjusts to growth gracefully.
- If data entries are data records, splits can change rids!
- **Key compression increases fanout, reduces height.**
- **Bulk loading can be much faster than repeated inserts for creating a B+ tree on a large data set.**
- **Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.**